


Advanced



操作系统
Operating System

Yu Chen

Tsinghua University

2015



Who ?

- **Lecturer: Chen Yu**
 - Associate Professor of Computer Science&Technology
 - yuchen@tsinghua.edu.cn
 - Interests:
 - **Operating Sytems**
- TA: Qixue Xiao, Junjie Mao,



Why Study OS?

- The Operating System (OS) I use has already been written, and I doubt it will be my job to write another one. For example, Windows, Linux.
- Haven't OS developers figured everything out already? What more is there to do?
- Why should I study this as a graduate student?



Why Study OS?

- The Operating System (OS) I use has already been written, and I doubt it will be my job to write another one. For example, Windows, Linux.
- Haven't OS developers figured everything out already? What more is there to do?
- Why should I study this as a graduate student?

OS is cool!





Why Study OS?

- The Operating System (OS) I use has already been written, and I doubt it will be my job to write another one. For example, Windows, Linux.
- Haven't OS developers figured everything out already? What more is there to do?
- Why should I study this as a graduate student?

OS is cool!

OS is important!





Why Study OS?


- The Operating System (OS) I use has already been written, and I doubt it will be my job to write another one. For example, Windows, Linux.
- Haven't OS developers figured everything out already? What more is there to do?
- Why should I study this as a graduate student?

OS is cool!

OS is important!

OS is challenging!

You want to be involved!





Objectives

- Gain experience in doing OS research
 - Know how to read/write papers/reports



Objectives

- Gain experience in doing OS research
 - Know how to read/write papers/reports
 - Know current OS hot topics



Objectives

- Gain experience in doing OS research
 - Know how to read/write papers/reports
 - Know current OS hot topics
 - Develop and test OS projects



Course Materials

- Lecture notes, Papers & Projects
- No required textbooks



Course Materials

- Lecture notes, Papers & Projects
- No required textbooks

Reference Books

- Wolfgang Mauerer, Professional Linux Kernel Architecture
- Uresh Vahalia, UNIX Internals-- The New Frontiers
- Daniel P. Bovet, Marco Cesati, Understanding the Linux Kernel,
- Mark E. Russinovich, David A. Solomon, Microsoft Windows? Internals
- Tanenbaum, Modern Operating Systems
- Andrew S. Tanenbaum, Distributed Operating Systems



Scheduling

Week	Lectures (includes some invited talks, Interviews)
1	Course overview + Speed+Security+Correct
2	OS Architecture
3	Virtual Machine
4	Building Real OS—uCore, Linux
5	Overview of OS Optimization for Multicore
6	Paper&Project: The Scalable Commutativity Rule & Commuter
7	Paper&Project: TCP/IP & FastSocket
8	Midterm Exam
9	Overview of OS Security
10	Paper&Project: Integer Security for System & Kint
11	Paper&Project: Lang, Library, Device Driver & RUST ,KLEE, PF-Miner
12	Overview of Correction on System
13	Paper&Project: Verified OS & SeL4
14	Paper&Project: Verified Component & JITK
15-16	Final Exam



Scheduling

Week	Lectures (includes some invited talks, Interviews)
1	Course overview + Speed+Security+Correct
2	OS Architecture
3	Virtual Machine
4	Building Real OS—uCore, Linux



Scheduling

Week	Lectures (includes some invited talks, Interviews)
5	Overview of OS Optimization for Multicore
6	Paper&Project: The Scalable Commutativity Rule & Commuter
7	Paper&Project: TCP/IP & FastSocket
8	Midterm Exam



Scheduling

Week	Lectures (includes some invited talks, Interviews)
9	Overview of OS Security
10	Paper&Project: Integer Security for System & Kint
11	Paper&Project: Lang, Library, Device Driver & RUST ,KLEE, PF-Miner



Scheduling

Week	Lectures (includes some invited talks, Interviews)
12	Overview of Correction on System
13	Paper&Project: Verified OS & SeL4
14	Paper&Project: Verified Component & JITK
15- 16	Final Exam



Reading Grading

- 6 Papers Summaries and critiques 60%
 - include report doc(4+ or 4000 words in A4 pages) & presentation (15+ pages)
 - Research Areas
 - OS/VMM Architectures
 - Performance/Multicore
 - Security/Find Bug
 - Correction/Verification



Reading Grading

- 6 Papers Summaries and critiques 60%
 - include report doc(4+ or 4000 words in A4 pages) & presentation (15+ pages)
 - Research Areas
 - OS/VMM Architectures
 - Performance/Multicore
 - Security/Find Bug
 - Correction/Verification
 - The **best papers** are from below conferences: SOSP, OSDI, EuroSys, RTSS, NSDI, FAST, USENIX ATC, VEE, MobiSys, PPOPP, SC, ICS, PLDI, SenSys, SIGMETRICS, ISCA, HPCA, HPDC, Micro, ASPLOS, IPDPS, SIGMOD, SIGCOMM, PerCom, WWW, Cluster, CCGrid, S&P, CCS, SECURITY, NDSS,... (1996~2015)



Reading Critiques

- Need to address the following:
 - Summary of major innovations
 - What the problems the paper mentioned?
 - How about the important related works/papers?
 - What are some intriguing aspects of the paper?



Reading Critiques

- Need to address the following:
 - Summary of major innovations
 - What the problems the paper mentioned?
 - How about the important related works/papers?
 - What are some intriguing aspects of the paper?
 - How to test/compare/analyze the results?
 - How can the research be improved?
 - If you write this paper, then how would you do?



Project Grading

- One Project 40%
 - didn't finish ucore/xv6/jos labs
 - Analyze/Testing/Improving ucore OS
 - <https://github.com/chyyuu/ucore>



Project Grading

- One Project 40%
 - didn't finish ucore/xv6/jos labs
 - Analyze/Testing/Improving ucore OS
 - https://github.com/chyyuu/ucore_lab
 - already finished ucore/xv6/jos labs
 - Do [one OS related project](#) and write the final report
 - research areas
 - There are a project lists
 - Discuss with me



Project Proposal

- Due on the 2nd~3rd week
- 2-page written proposal
 - team members (≤ 3)
 - Motivation
 - The state-of-the-art Methodology
 - Expected results
 - Timeline
 - Division of labor among teams
 - Some references



Project Midterm/Final Exam

- Needed
 - 10-15 minutes Presentation
 - Midterm
 - 3~10 page written paper/report (double column, single-space, 10-pt font)
 - Final
 - 5~15 page written paper/ project report (double column, single-space, 10-pt font)



Other Choose

- I can not understand the paper, my coding ability is poor
 - Need to talk with me ASAP



OS Overview



History of OS: Change!

		1980	2009	Factor
Speed	CPU	1 MIPS	77,000 MIPS	7.7×10^5
	Memory	500 ns	0.9 ns	1.8×10^3
	Disk	18 ms	4 ms	2.2×10^1
	Network	300 b/sec	10 Gb/sec	3.3×10^7
Capacity	Memory	64 KB	8 GB	1.3×10^5
	Disk	1 MB	2 TB	2.0×10^6
Cost	Per MIP	\$100K/MIP	\$.007899	1.0×10^7
Other	Address bits	8	64	8
	Users/machine	10s	0.1	1.0×10^{-2}



Changing Roles of the OS

- What OS does depends on
 - available hardware and software
 - changing uses of machines
 - changing expectations of users



OS Concepts

- Single-Machine OS/VMM
 - Memory management
 - Process management
 - Synchronization
 - File systems and device support
- Security
- Correction
- Distributed System



Single-Machine OS/VMM

Purposes

- Clean virtual machine
- Hardware independence
- Resource sharing and management
- Persistent Data Storage
- Protection
- Real time support
- Parallelism



Single-Machine OS/VMM

Purposes

- Clean virtual machine
- Hardware independence
- Resource sharing and management
- Persistent Data Storage
- Protection
- Real time support
- Parallelism

Strategy

- How do we organize the OS effectively for development, evolution, performance, and security?
- How do we use multi-processor machines effectively?



Memory Management

Purposes

- Virtual memory: provides the illusion of infinite physical memory
- Swapping: moves processes to disk as necessary
- Paging: allows processes to run with only the active pages in memory
- Buffer Cache: speedup the IO access



Memory Management

Purposes

- Virtual memory: provides the illusion of infinite physical memory
- Swapping: moves processes to disk as necessary
- Paging: allows processes to run with only the active pages in memory
- Buffer Cache: speedup the IO access

Strategy

- How do we coordinate machines to share memory?
- How can we simplify memory management as memory becomes abundant?



Process Management

- Thread:
- Address space:
- Process:



Process Management

- Thread: A sequential execution stream
- Address space: Chunks of memory and everything needed to run a program
- Process: An address space + thread(s)



Process Management

Purposes

- Thread: A sequential execution stream
- Address space: Chunks of memory and everything needed to run a program
- Process: An address space + thread(s)

Strategy

- How do processes communicate and share states efficiently and securely on the same machine?
- How do we improve the computing process model?



Process Scheduling

Purposes

- Provides the illusion of multiple processes running at the same time on a single processor
- Context switching: changing the attention of the processor
 - Involves saving and restoring states
 - Necessary to cross kernel boundary



Process Scheduling

Purposes

- Provides the illusion of multiple processes running at the same time on a single processor
- Context switching: changing the attention of the processor
 - Involves saving and restoring states
 - Necessary to cross kernel boundary

Strategy

- How do we achieve fairness, high throughput, and responsiveness at the same time?
- How do we reduce or avoid the cost of context switching?



Memory Management

Purposes

- Virtual memory: provides the illusion of infinite physical memory
- Swapping: moves processes to disk as necessary
- Paging: allows processes to run with only the active pages in memory
- Buffer Cache: speedup the IO access



Memory Management

Purposes

- Virtual memory: provides the illusion of infinite physical memory
- Swapping: moves processes to disk as necessary
- Paging: allows processes to run with only the active pages in memory
- Buffer Cache: speedup the IO access

Strategy

- How do we coordinate machines to share memory?
- How can we simplify memory management as memory becomes abundant?



Synchronization

Purposes

- Provides correct execution or coordinating threads in the face of arbitrary context switching



Synchronization

Purposes

- Provides correct execution or coordinating threads in the face of arbitrary context switching

Strategy

- Atomic actions: all or nothing
- Mutual exclusion: one thread in the critical section at a time
- Semaphores: atomic, counter-based locks
- Avoid Deadlock: circular waiting on resources



File Systems

Purposes

- File: data + attributes
- File system services:
 - Organization
 - Naming
 - Access
 - Synchronization
 - Protection and security



File Systems

Purposes

- File: data + attributes
- File system services:
 - Organization
 - Naming
 - Access
 - Synchronization
 - Protection and security

Strategy

- How do we make different file systems work together, even across machines?
- How do we provide consistency, availability, and reliability to copies of a file across multiple machines?
- How do we handle very large data sets?



I/O Device

Purposes

- I/O devices tend to be a lot slower than memory speed
- Caching: stores extra data in memory in hope of near-term reuse



I/O Device

Purposes

- I/O devices tend to be a lot slower than memory speed
- Caching: stores extra data in memory in hope of near-term reuse

Strategy

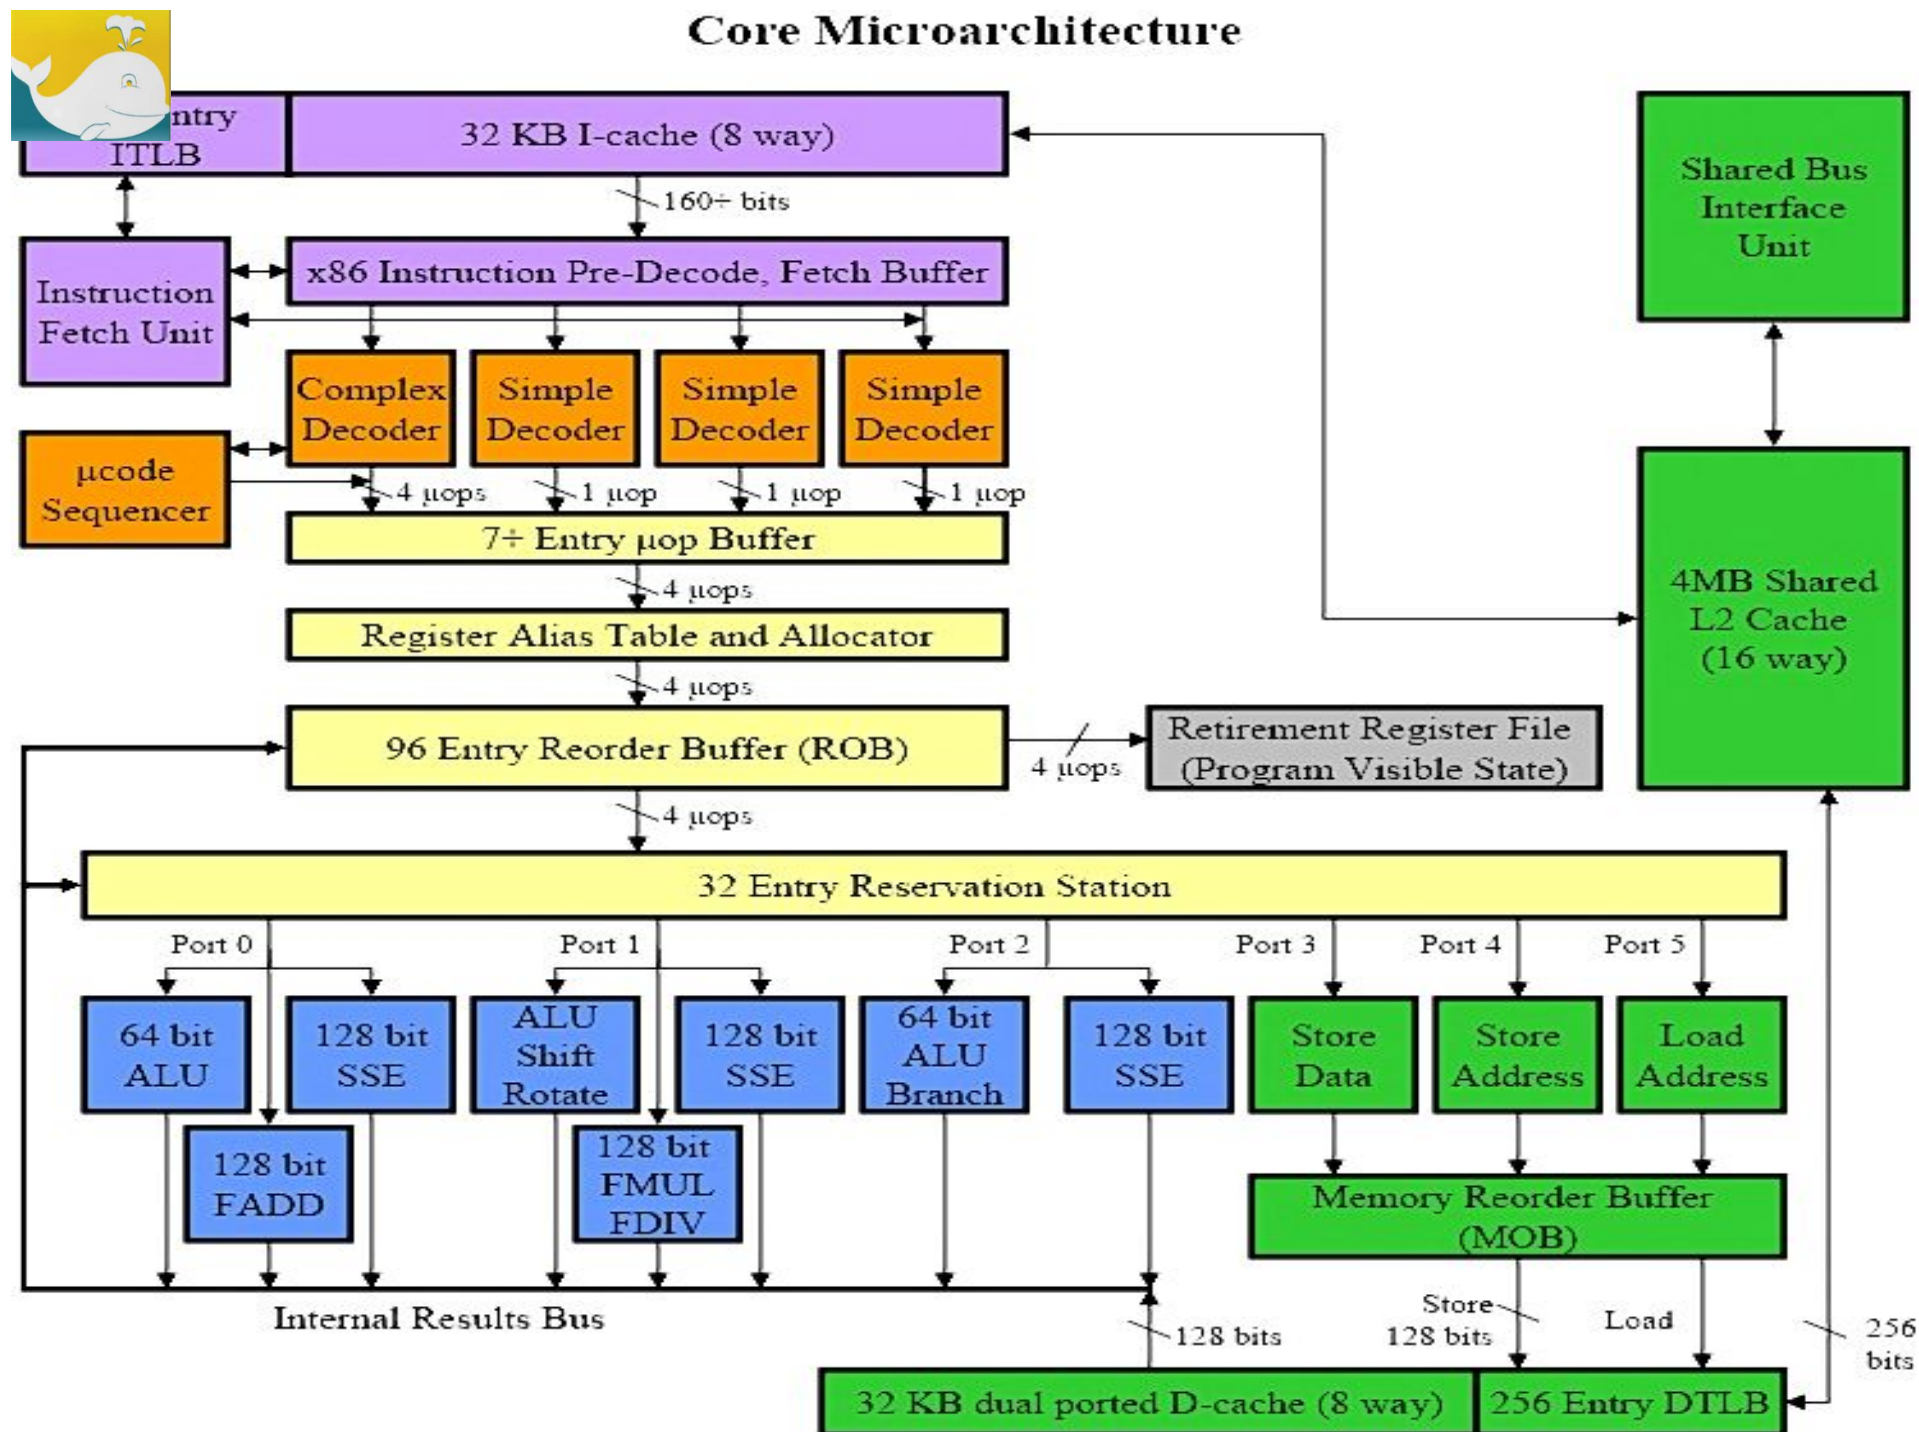
- How do we coordinate the memory resources across machines to enhance performance?
- How do we handle new devices with new characteristics?



Tendency

- MultiCore
- Security
- Correction

Core Microarchitecture





Multi-Core Challenges

- Today's Commodity Multi-Cores




Fig. 1. The architecture of Intel's multi-core platform used in experiments. Squares with *C* tag represent cores, squares with *L1* tag represent L1 instruction cache and L1 data cache, rectangles with *L2* represent L2 unified cache, rectangle with *DRAM* represents memory, circle with *MCH* represents memory controller hub, lines with arrow represent directly accessible relationship and rectangles with *CPU* represent chips. Each cache line is 64 bytes, both L1 instruction cache and L1 data cache are 32K bytes 8-way set associative, and the L2 cache is 4M bytes 16-way set associative.

Memory References	Latency
memory read	212 cycles
L1 data cache hit	3 cycles
L2 unified cache hit	14 cycles
local cache-to-cache transfer	46 cycles
remote cache-to-cache transfer	50 cycles

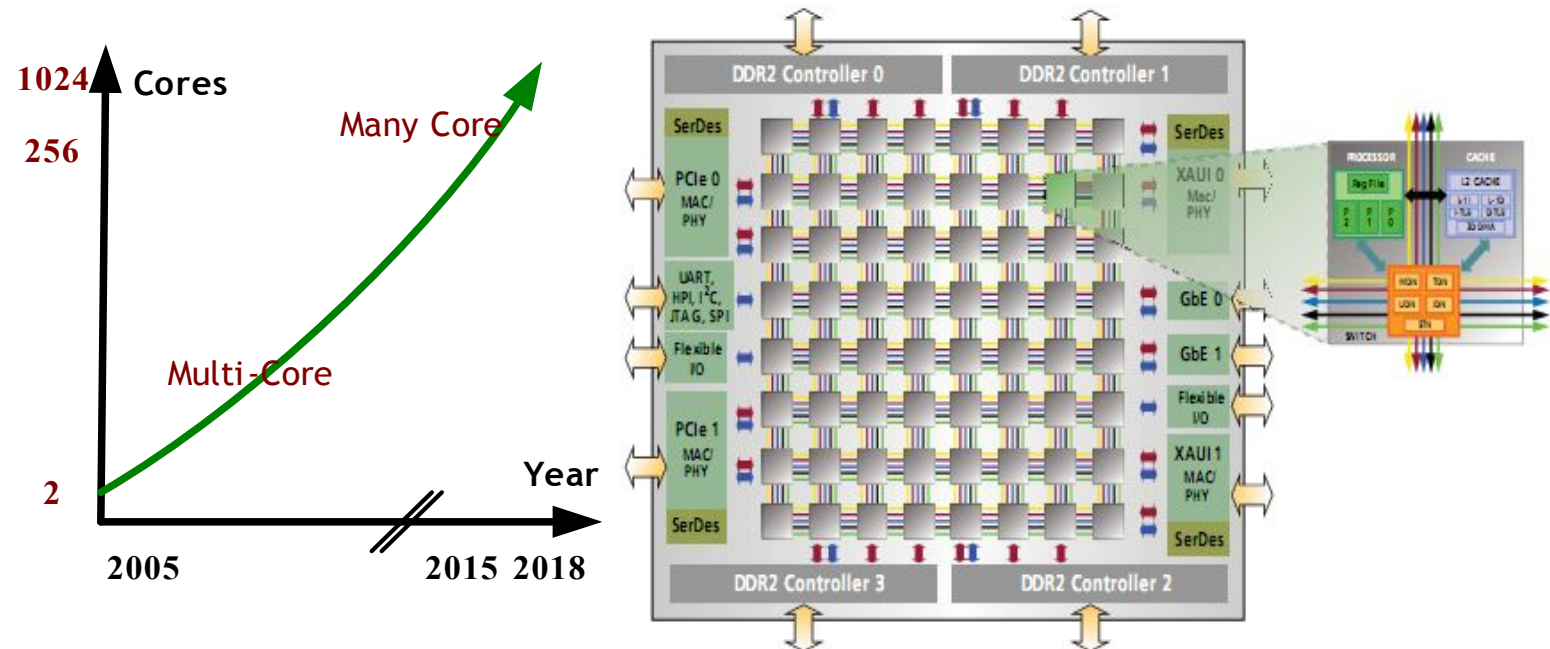
Applications should avoid the memory wall

Cache-to-cache transfer should be avoided



Multi-Core Challenges

- CMP V.S. SMP
 - **More cores can be integrated**
 - SMP: Low(2CPU),Middle(4~8CPU),High(>16CPU)
 - CMP: 4~8 cores systems, **1000+ cores (<10years)**
E.g. Intel's 80 cores chip & Tiler's 64 core chip





Multi-Core Challenge

- However, can operating systems and applications use these cores effectively?





Fig 2. Performance bound by bottlenecks





OS supporting Multicore

- Linux
- Solaris
- FreeBSD
- Windows
- VxWorks





Benchmarks





Some Conclusions

- No system scales clearly better than another in all aspects for micro-benchmark test
- Linux and Solaris are competitive in application benchmark test, FreeBSD loses both in performance and scalability
- Kernel synchronizations protecting the shared data structure are the main bottlenecks on multi-core platform



Tendency

- MultiCore
- Security
- Correction



History of Security Problem

- Originally, there was no security problem
- Later, there was a problem, but nobody cared
- Now, there are increasing problems, and people are beginning to care
 - Automation
 - Action at a distance
 - Technique propagation



Threat Analysis

- What are we trying to protect? (and why?)
- What are the vulnerabilities of those assets?
- Who might exploit a vulnerability?
 - Either on purpose or by accident



Threat Analysis

- What are we trying to protect? (and why?)
- What are the vulnerabilities of those assets?
- Who might exploit a vulnerability?
 - Either on purpose or by accident
- How can we prevent a specific threat?
- How much is it worth to us to prevent it?
- How much will it cost to prevent it?



The Core Technical Problem

- Controlling access to machine and data resources
- Controlling the way *access rights* are passed from holder to holder
 - person to person
 - program to program
- Preventing maliciousness and errors from subverting the controls



Access Rights

- In general case, need triplet for every possible combination of *right*, protected *asset*, and potential user
 ↪ (U, A, R)
- ...and some entity must be responsible for checking and enforcing any limitation...
- The 3-D matrix is hard to manage...We need a simpler approach!



ACL&Capabilities

	<i>file1</i>	<i>file 2</i>	<i>file 3</i>	<i>device</i>	<i>domain</i>
<i>User/Domain 1</i>	<i>r</i>	<i>rx</i>	<i>rw</i>	—	<i>enter</i>
<i>User/Domain 2</i>	<i>r</i>	<i>x</i>	<i>rx</i>	<i>rw</i>	—
<i>User/Domain 3</i>	<i>rw</i>	—	—	—	—
...					

- Columns are *access control lists* (ACLs)
 - Associated with each object
- Rows are *capabilities*
 - Associated with each user or each domain




What about efficiency?

- At run-time...
 - What does the OS know about the user?
 - What does the OS know about the resources?
- What is the cost of checking and enforcing?
 - Access to the data
 - Cost of searching for a match



Current Status (1)



- 对当前Android漏洞的理解
 - Sematic Vulnerability 越来越多
 - 数据泄漏漏洞的威胁越来越大





Current Status (2)

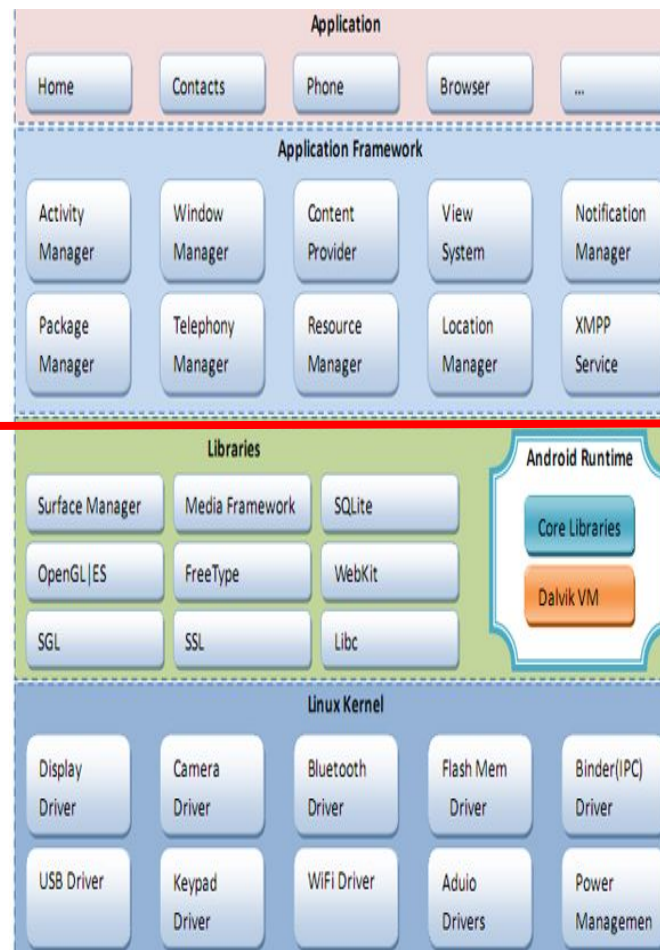
- 对当前Linux Kernel漏洞的理解
 - Linux漏洞有扩大化的趋势
 - 但发现Linux漏洞难度加大





System Security Technology

安全分析工具



权限模型
Binder -Intent机制
Content Provider机制
应用签名验证机制
用户隐私保护机制

SQLite数据库机制
SSL机制

内核级整数溢出
内核级内存溢出
内核数据流分析
内核驱动缺陷分析

系统安全机制

符号执行

数据流分析

动态执行分析

静态控制流分析

Model

Check

核心分析技术



Tendency

- MultiCore
- Security
- Correction




Standard Motivating Slide for Verification



Ariane 5



Mars Polar
Lander



Mars climate
orbiter



Photo from edmunds.com

Toyota recalled its 160,000 Prius cars in Oct 2005, because of bugs in the software controlling the hybrid gas-electric engine system...

Report: Prius hybrids dogged by software woes - May. 16, 200...

File Edit View Go Bookmarks Tools Help

http://money.cnn.cor

Search Postings Getting Started Latest Headlines Overview (Java 2 Platf...

The Internet home of: FORTUNE Money BUSINESS 2.0 FORTUNE SMALL BUSINESS Subscribe to

CNNMoney.com Autos In partnership with edmunds.com

HOME NEWS MARKETS TECHNOLOGY JOBS & ECONOMY PERSONAL FIN

New Cars Certified Cars Used Cars Car Reviews Ti

SAVE | EMAIL | PRINT | Re

Prius hybrids dogged by software

Report: Internal computer woes reportedly cause autos to stall or shut down at highway speeds.

May 16, 2005: 1:15 PM EDT

NEW YORK (CNN/Money) - A software problem is causing some Toyota Prius gas-electric hybrid cars to stall or shut down while driving at highway speeds, according to a published report.

The *Wall Street Journal* reports that the problem involves Priuses from the 2004 model year and some early 2005 models.

RESEARCH A CAR
Get invoice and market prices, specs, reviews and photos

- Sport • Sedans
- SUVs • Luxury

Pick Category

Go

Check latest incentives

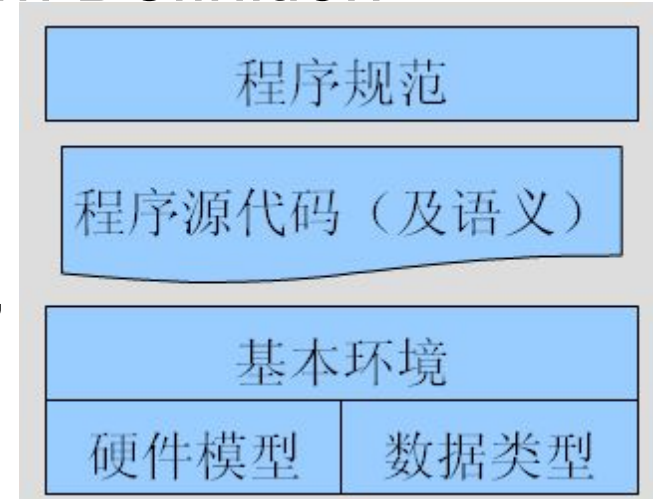
Done



Basic Idea

Certified Software: Problem Definition

- **Hardware**
 - processors, memory, storage, devices,
- **Software**
 - bootloader, device drivers, OS, runtime, applications, ...
- Need a mathematical **proof** showing that as long as the **hardware** works, the **software** always work according to its **specification**





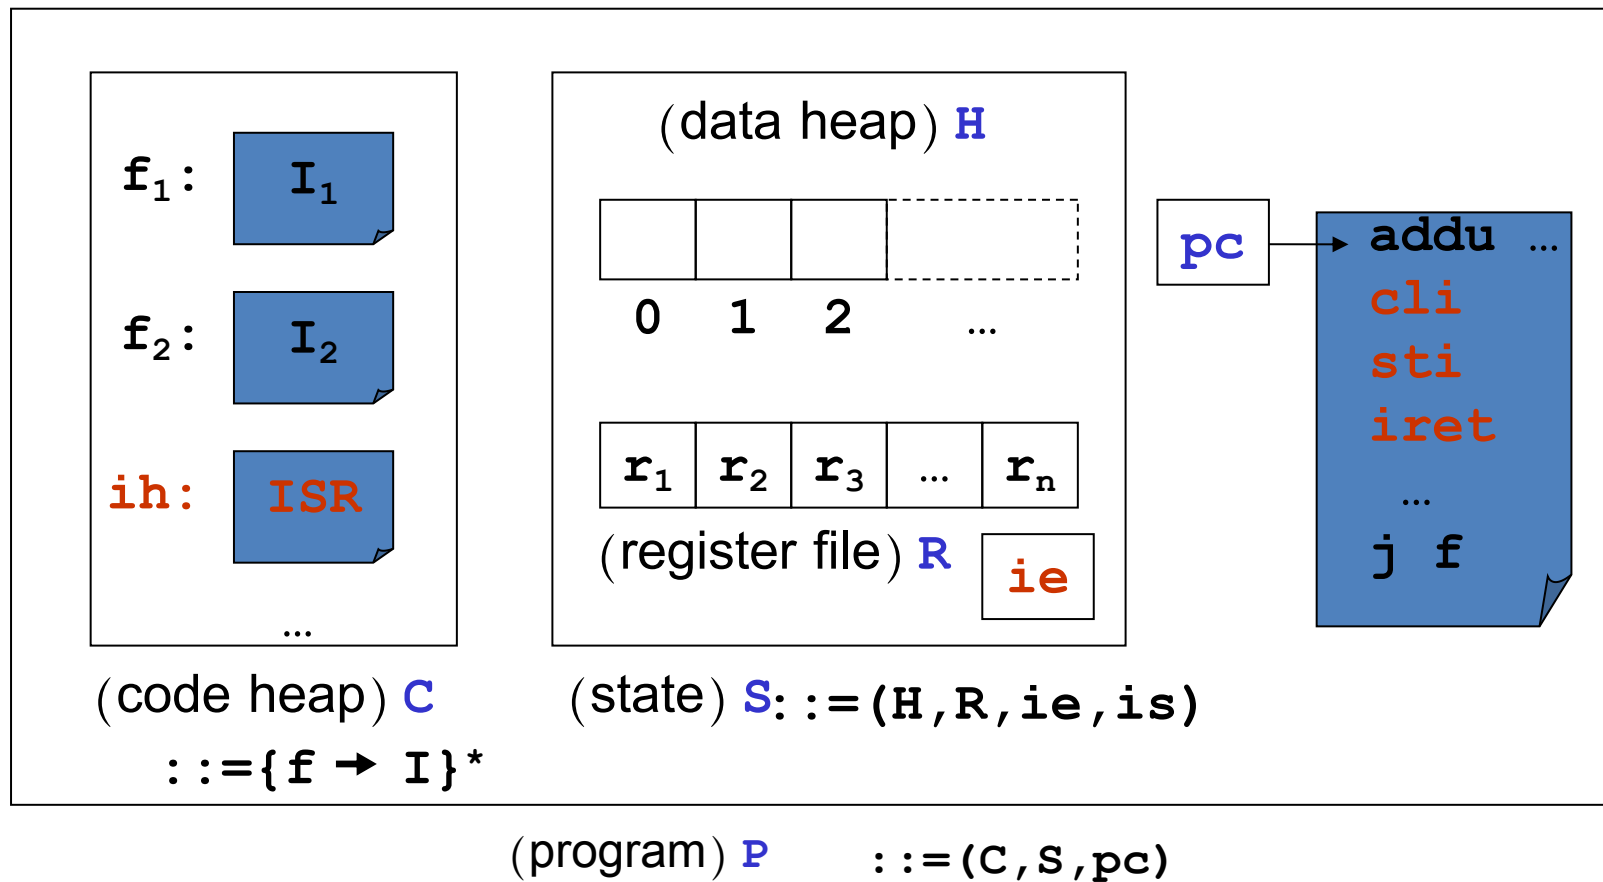
Challenges

- Components come from different sources
 - Manually written assembly
 - C/C++
 - Type safe languages
 - Java, C# ...
- Many different features
 - Code loading
 - Control abstractions
 - jmp (goto)
 - functions
 - exceptions
 - threads
 - interrupts



AIM : The Machine

Abstract Interrupt Machine





Challenges

- Components come from different sources
 - Manually written assembly
 - C/C++
 - Type safe languages
 - Java, C# ...
- Many different features
 - Memory update
 - type-preserving update
 - type-changing update
 - pointer arithmetic
 - Device drivers and I/O
 - ...



Some Conclusions

- AIM machine
 - low-level
 - can implement interrupt handlers and thread libraries
- A program logic
 - following local reasoning in separation logic
 - modeling `cli/sti`, `switch`, `block/unblock` in terms of memory ownership transfer
 - can certify different implementation of locks and C.V.s
- mini-OS proof
 - Interrupt handler, scheduler and others



Thanks

- Some materials are from Andy Wang
- CS-502 Operating Systems from WPI
- Compiler/Program Research Group in TH
-