

Reference Count Imbalance Detection without Semantic Assumptions

Junjie Mao

Tsinghua University

5/26/2015

Outline

- Introduction
- Related Work on Detecting Reference Count Bugs
- Inconsistency Analysis
- Implementation
- Preliminary Evaluation
- Future Work

Introduction

- Reference counts (refcount) are widely used in OS kernels from servers to mobile devices, ...
 - memory management
 - power management
 - ... [see backup pages]
- ... but refcount imbalances can lead to critical system failures.
 - out of memory, dangling pointers, ...
 - send requests to devices in low-power state
 - etc.

Related Work

- Related works have proposed algorithms for detecting possible refcount imbalances on affine programs with shallow alias [1,2], ...

```
int foo(i, x1, x2) {  
    if(i > 0) {  
        x3 = x1;  
    } else {  
        x3 = x2;  
    }  
    if(i > 0) {  
        x3->refcount ++;  
        return 1;  
    }  
    return 0;  
}
```

affine



```
global x1, x2  
  
foo() {  
    if ? {  
        x3 = x1;  
    } else {  
        x3 = x2;  
    }  
    if ? {  
        x3->refcount ++;  
    }  
}
```

[1] A. Lal and G. Ramalingam. Reference count analysis with shallow aliasing. Information Processing Letters, 111(2):57–63, 2010

[2] S. Li and G. Tan. Finding reference-counting errors in python/C programs with affine analysis. In ECOOP 2014

Related Work (cont.)

- Related works have proposed algorithms for detecting possible refcount imbalances on affine programs with shallow alias, ...
 - ... with the assumption that refcounts used should always be balanced.

```
foo() {  
    set all refcounts to 0  
    .....  
    check all refcounts are still 0  
    (or other assumed value)  
}
```

Related Work (cont.)

- The techniques above do not apply to an OS kernel, however.
 - The entry of an OS kernel is expected to be “noreturn”
 - Functions in an OS kernel may be designed to have non-balanced refcount operations
- How can we know a refcount imbalance is bound to happen **without the semantic assumption?**

Related Work (cont.)

- In [3] and [4], inconsistency among beliefs from different code segments are used to check potential errors statically.


```
bar(x) {  
    if (x != NULL) *x;  
    ...  
    *x;  
    ...  
}
```

```
lock l;  
int a, b;  
void foo0 {  
    lock(l); a = a + b; unlock(l);  
    b = b + 1;  
}  
void bar() {  
    lock(l); a = a + 1; unlock(l);  
}  
void baz0 {  
    a = a + 1;  
    unlock(l); b = b - 1; a = a / 5;  
}
```

[3] D. R. Engler, D. Y. Chen, and A. Chou. Bugs as inconsistent behavior: A general approach to inferring errors in systems code. In SOSP, pages 57–72, 2001.

[4] I. Dillig, T. Dillig, and A. Aiken. Static error detection using semantic inconsistency inference. In ACM SIGPLAN Notices, volume 42, pages 435–445, June 2007.

Inconsistency Analysis



```
void foo(x1, x2) {  
    if ? {  
        x3 = x1;  
    } else {  
        x3 = x2;  
    }  
    if ? {  
        x3->refcount ++;  
    }  
}
```

a specific path

```
x2->refcount --;  
// x2->refcount should be 0  
// after the decrement
```

- In the affine program, a function with different refcount operations on different path will always raise refcount imbalance.
- However, this too imprecise to be practical.

Inconsistency Analysis (cont.)

```
int foo(i, x1, x2) {  
    if(i > 0) {  
        x3 = x1;  
    } else {  
        x3 = x2;  
    }  
    if(i > 0) {  
        x3->refcount ++;  
        return 1;  
    }  
    return 0;  
}
```

- Is foo() really buggy? We are still no sure, because:
 - If $i > 0$, the refcount of x1 is incremented.
 - If $i \leq 0$, no refcount is changed.
 - Callers of foo() can determine how foo() will affect refcounts based on the actual parameter passing to foo() or the return value of foo().
 - It is the caller's responsibility to handle both cases properly.

Inconsistency Analysis (cont.)

- Callers to `foo()` may determine the refcount operations of `foo()` by:
 - parameters passed to `foo()`
 - the return value of `foo()`
 - side effects of `foo()` on global data structures
- In this work we focus on the first two cases which assumes all functions are pure except their refcount operations.

Inconsistency Analysis (cont.)

C construct	affine translation in previous work	our affine translation
int foo(a1, ..., ak) {...}	foo() {...}	int foo(a1, ..., ak) {...}
if (i > 0) {...} else {...}	if ? {...} else {...}	if (i > 0) {...} else {...}
if (bar(...)) {...} else {...}	if ? {...} else {...}	if ? {...} else {...}
return v;	(none)	return v;
x = f(x1, ..., xk)	f()	x = f(x1, ..., xk)
x = y	x = y	x = y

- We adopt a summary-based inter-procedural function analysis to implement the detection process. Functions are visited in reverse topological order.

Function/ summary Based IPA

- We summarize a function by recording
 - modified refcounts
 - under what path conditions are these refcounts modified, and
 - how the refcounts are modified.
- This form of summary allows us to carry out optimiations on path enumeration which is discussed later.

```
int foo(i, x1, x2) {  
    if(i > 0) { x3 = x1; }  
    else { x3 = x2; }  
    if(i > 0) {  
        x3->refcount ++;  
        return 1;  
    }  
    return 0;  
}
```

summarize →

Summary		
refcount	cond / op	
x1	[i]>0 ^ [1]* =	+1
x1	[i]<=0 ^ [0] =	0

* [] represents the return value

Function/ summary Based IPA (cont.)

- Summaries are applied in call sites based on the condition of actual argument and return values

Summary of foo()

refcount	cond / op
x1	$[i] > 0 \wedge [1]^* = +1$
x1	$[i] \leq 0 \wedge [0] = 0$

```
int bar(k, x1, x2) {  
    ret = foo(k, x1, x2);  
    if (ret == 0) {  
        x1->refcount ++;  
    }  
}
```

- path 1: $[k] \leq 0 \wedge \text{ret} = 0$
 - $x1$
- path 2: $[k] > 0 \wedge \text{ret} = 1$
 - $x1 + 1$

Bug Reporting

- The summaries are calculated in the reverse topological order of the call graph.
- Bugs are reported in the following case

refcount	cond / op
x1	$[i] > 0 \wedge [1] = 0$
x1	$[i] > 0 \wedge [1] = +1$

- Note that the following case may not be a bug

refcount	cond / op
x1	$[i] > 0 \quad 0$
x1	$[i] = 0 \quad +1$


applied to a path whose condition is $[i] \geq 0$

Implementation

- The IPA analysis is implemented based on the Static Single Assignment (SSA) form of Control Flow Graph (CFG)
- For each function, all *trails* in the CFG are enumerated and function summaries are applied to the call sites.
 - A *trail* in a CFG is a path from the entry point to the exit point with no repeated edges.


Analysis on a Single Path

- An entry in the summary may correspond to multiple paths in the function.
 - Different paths can have the same path condition and call sites
- A path in the function may also correspond to multiple entries in the summary.



Implementation

- However, the number of paths in a CFG is exponential to the number of branches in general.
 - We have met a function with over 1 billion paths in the Linux kernel source
- Before path enumeration, we simplify the CFG by removing *ineffective* vertices (i.e. basic blocks).




Remove Ineffective Vertices

- We first determine the set of *effective* instructions in a CFG. An instruction is *effective* if one of the following holds
 - It is a call to a function with refcount operations (this instruction defines a variable holding the return value at the same time)
 - It defines a variable which is used in an effective instruction
 - It uses a variable defined by an effective instruction to define another variable
 - It is a comparison on an effective instruction which is used in a branch condition
- Vertices which covers no instruction in the effective instruction set are ineffective, and are deleted on after another during the graph simplification.

Preserving Indirect Data Dependency

- Two sets are attached to each edge
 - pred: set of predicates that held by walking from its source vertex to the target vertex
 - edge: set of edges in the original CFG that will always be passed when walking from its source vertex to the target
- Initialization




$e_1.\text{pred}: \{i > 0\}$
 $e_1.\text{edge}: \{e_1\}$


$e_2.\text{pred}: \{i \leq 0\}$
 $e_2.\text{edge}: \{e_2\}$

Preserving Indirect Data Dependency

- For newly created edges during the simplification of the CFG:



$$\begin{aligned} e'.pred &= e_1 \cup e_2 \\ e'.edge &= e_1 \cup e_2 \end{aligned}$$



$$\begin{aligned} e'.pred &= e_1 \cap e_2 \\ e'.edge &= e_1 \cap e_2 \end{aligned}$$

Preserving Indirect Data Dependency

- A trail is represented by a sequence of edges:

$$e_1, e_2, \dots, e_k$$

- Path conditions of a trail (written as PC) in the simplified CFG is:

$$AP = \bigcup_{i=1}^k e_i.path$$
$$PC = \{p \mid p \in AP \wedge \neg p \notin AP\}$$

- For a ϕ variable in v_k whose in edge is e_k in the trail, the value of this variable is then determined by $e_k.edge$

Preliminary Evaluation

- Common error types
 - Missing resource release in error-handling path
 - Jumping to the wrong exception handler
 - Wrong assumption on function behavior

Preliminary Evaluation: Example

```
static void g2d_dma_start(struct g2d_data *g2d,  
                          struct g2d_runqueue_node *runqueue_node)  
{  
    struct g2d_cmdlist_node *node =  
        list_first_entry(&runqueue_node->run_cmdlist,  
                        struct g2d_cmdlist_node, list);  
  
    int ret;  
  
    ret = pm_runtime_get_sync(g2d->dev);  
    if (ret < 0)  
        return;  
  
    writel_relaxed(node->dma_addr, g2d->regs + G2D_DMA_SFR_BASE_ADDR);  
    writel_relaxed(G2D_DMA_START, g2d->regs + G2D_DMA_COMMAND);  
}
```

Summary

refcount	cond / op	
g2d->dev- >power.usage_count	True	+1

Future Work

- Symbolic-execution based summary calculation
- Consideration on common side effects
- More precise analysis on loops

Thanks!

Backup

linux kernel 2.6.5 引入kref

```
struct foo {  
    ...  
    struct kref kref;  
    ...  
}
```

```
struct foo *foo;  
foo = kmalloc(sizeof(*foo),GFP_KERNEL);  
kref_init(&foo->kref, oo_release);
```

```
void foo_release(struct kref *kref)  
{  
    struct foo *foo;  
    foo = container_of(kref, struct foo,kref);  
    kfree(foo);  
}
```

```
struct kref {  
    atomic_t refcount;  
    void (*release)(struct kref *kref);  
};
```

Backup

linux kernel 2.6.5 引入kref

```
void kref_init(struct kref *kref, void (*release)
(struct kref *kref))
{  WARN_ON(release == NULL);
   atomic_set(&kref->refcount,1);
   kref->release = release;
}
```

```
struct kref *kref_get(struct kref *kref)
{  WARN_ON(!atomic_read(&kref->refcount));
   atomic_inc(&kref->refcount);
   return kref;
}
```

```
void kref_put(struct kref *kref)
{  if (atomic_dec_and_test(&kref->refcount))
    kref->release(kref);
}
```